

# Package ‘ngram’

March 24, 2017

**Type** Package

**Title** Fast n-Gram 'Tokenization'

**Version** 3.0.3

**Description** An n-gram is a sequence of n ``words" taken, in order, from a body of text. This is a collection of utilities for creating, displaying, summarizing, and ``babbling" n-grams. The 'tokenization' and ``babbling" are handled by very efficient C code, which can even be built as its own standalone library. The babbler is a simple Markov chain. The package also offers a vignette with complete example 'workflows' and information about the utilities offered in the package.

**License** BSD 2-clause License + file LICENSE

**Depends** R (>= 3.0.0)

**Imports** methods

**LazyData** yes

**LazyLoad** yes

**NeedsCompilation** yes

**ByteCompile** yes

**Maintainer** Drew Schmidt <wrathematics@gmail.com>

**URL** <https://github.com/wrathematics/ngram>

**BugReports** <https://github.com/wrathematics/ngram/issues>

**Collate** 'ngram.r' 'babble.r' 'checks.r' 'concatenate.r' 'getseed.r'  
'getters.r' 'multiread.r' 'ngram-package.R' 'ngram\_asweka.r'  
'phrasetable.r' 'preprocess.r' 'print.r' 'rcorpus.r'  
'splitter.r' 'string.summary.r' 'wordcount.r'

**RoxygenNote** 5.0.1

**Author** Drew Schmidt [aut, cre],  
Christian Heckendorf [aut]

**Repository** CRAN

**Date/Publication** 2017-03-24 05:37:04 UTC

R topics documented:

ngram-package . . . . .	2
babble . . . . .	3
concatenate . . . . .	4
getseed . . . . .	5
getters . . . . .	5
multiread . . . . .	6
ngram-class . . . . .	7
ngram-print . . . . .	8
phrasetable . . . . .	9
preprocess . . . . .	9
rcorpus . . . . .	10
splitter . . . . .	11
string.summary . . . . .	12
Tokenize . . . . .	12
Tokenize-AsWeka . . . . .	13
wordcount . . . . .	14
<b>Index</b>	<b>16</b>

---

ngram-package	<i>ngram: An n-gram Babbl</i>
---------------	-------------------------------

---

Description

An n-gram is a sequence of n "words" taken from a body of text. This package offers utilities for creating, displaying, summarizing, and "babbling" n-grams. The tokenization and "babbling" are handled by very efficient C code, which can even be build as its own standalone library. The babbler is a simple Markov chain.

Details

The ngram package is distributed under the permissive 2-clause BSD license. If you find the code here useful, please let us know and/or cite the package, whatever is appropriate.

The package has its own PRNG; we use an implementation of MT1997 for all non-deterministic choices.

Author(s)

Drew Schmidt <wrathematics AT gmail.com>, Christian Heckendorf.

---

babble*ngram Babbler*

---

## Description

The babbler uses its own internal PRNG (i.e., not R's), so seeds cannot be managed as with R's seeds. The generator is an implementation of MT19937.

At this time, we note that the seed may not guarantee the same results across machines. Currently only Solaris produces different values from mainstream platforms (Windows, Mac, Linux, FreeBSD), but potentially others could as well.

## Usage

```
babble(ng, genlen = 150, seed = getseed())
```

```
## S4 method for signature 'ngram'  
babble(ng, genlen = 150, seed = getseed())
```

## Arguments

ng	An ngram object.
genlen	Generated length, i.e., the number of words to babble.
seed	Seed for the random number generator.

## Details

A markov chain babbler.

## See Also

[ngram](#), [getseed](#)

## Examples

```
library(ngram)  
  
str <- "A B A C A B B"  
ng <- ngram(str)  
babble(ng, genlen=5, seed=1234)
```

---

concatenate	<i>Concatenate</i>
-------------	--------------------

---

## Description

A quick utility for concatenating strings together. This is handy because if you want to generate the n-grams for several different texts, you must first put them into a single string unless the text is composed of sentences that should not be joined.

## Usage

```
concatenate(..., collapse = " ", rm.space = FALSE)
```

## Arguments

...	Input text(s).
collapse	A character to separate the input strings if a vector of strings is supplied; otherwise this does nothing.
rm.space	logical; determines if spaces should be removed from the final string.

## Value

A string.

## See Also

[preprocess](#)

## Examples

```
library(ngram)

words <- c("a", "b", "c")
wordcount(words)
str <- concatenate(words)
wordcount(str)
```

---

`getseed`*getseed*

---

**Description**

A seed generator for use with the ngram package.

**Usage**

```
getseed()
```

**Details**

Uses a 96-bit hash of the current process id, time, and a random uniform value from R's random generator.

**See Also**

[babble](#)

---

`getters`*ngram Getters*

---

**Description**

Some simple "getters" for ngram objects. Necessary since the internal representation is not a native R object.

**Usage**

```
get.ngrams(ng)
```

```
## S4 method for signature 'ngram'  
get.ngrams(ng)
```

```
get.string(ng)
```

```
## S4 method for signature 'ngram'  
get.string(ng)
```

```
get.nextwords(ng)
```

```
## S4 method for signature 'ngram'  
get.nextwords(ng)
```

Arguments

ng                      An ngram object.

Details

get.ngrams() returns an R vector of all n-grams.  
get.nextwords() does nothing at the moment; it will be implemented in future releases.  
getnstring() recovers the input string as an R string.

See Also

[ngram-class](#), [ngram](#)

Examples

```
library(ngram)

str <- "A B A C A B B"
ng <- ngram(str)
get.ngrams(ng)
```

---

multiread	<i>Multiread</i>
-----------	------------------

---

Description

Read in a collection of text files.

Usage

```
multiread(path = ".", extension = "txt", recursive = FALSE,
  ignore.case = FALSE, prune.empty = TRUE, pathnames = TRUE)
```

Arguments

path	The base file path to search.
extension	An extension or the "*" wildcard (for everything). For example, to read in files ending .txt, you could specify extension="txt". For the purposes of this function, each of *.txt, *txt, .txt, and txt are treated the same.
recursive	Logical; should the search include all subdirectories?
ignore.case	Logical; should case be ignored in the extension? For example, if TRUE, then .r and .R files are treated the same.
prune.empty	Logical; should empty files be removed from the returned list?
pathnames	Logical; should the full path be included in the names of the returned list.

## Details

The extension argument is not a general regular expression pattern, but a simplified pattern. For example, the pattern `*.txt` is really equivalent to `*[.]txt$` as a regular expression. If you need more complicated patterns, you should directly use the `dir()` function.

## Value

A named list of strings, where the names are the file names.

## Examples

```
## Not run:
path <- system.file(package="ngram")

### Read all files in the base path
multiread(path, extension="*")

### Read all .r/.R files recursively (warning: lots of text)
multiread(path, extension="r", recursive=TRUE, ignore.case=TRUE)

## End(Not run)
```

---

ngram-class

*Class ngram*


---

## Description

An n-gram is an ordered sequence of n "words" taken from a body of "text". The terms "words" and "text" can easily be interpreted literally, or with a more loose interpretation.

## Details

For example, consider the sequence "A B A C A B B". If we examine the 2-grams (or bigrams) of this sequence, they are

A B, B A, A C, C A, A B, B B

or without repetition:

A B, B A, A C, C A, B B

That is, we take the input string and group the "words" 2 at a time (because  $n=2$ ). Notice that the number of n-grams and the number of words are not obviously related; counting repetition, the number of n-grams is equal to

$$nwords - n + 1$$

Bounds ignoring repetition are highly dependent on the input. A correct but useless bound is

$$\#ngrams = nwords - (\#repeats - 1) - (n - 1)$$

An ngram object is an S4 class container that stores some basic summary information (e.g.,  $n$ ), and several external pointers. For information on how to construct an ngram object, see [ngram](#).

**Slots**

str\_ptr A pointer to a copy of the original input string.

strlen The length of the string.

n The eponymous 'n' as in 'n-gram'.

ngl\_ptr A pointer to the processed list of n-grams.

ngsize The length of the ngram list, or in other words, the number of unique n-grams in the input string.

sl\_ptr A pointer to the list of words from the input string.

**See Also**

[Tokenize](#)

---

ngram-print

*ngram printing*


---

**Description**

Print methods.

**Usage**

```
## S4 method for signature 'ngram'
print(x, output = "summary")
```

```
## S4 method for signature 'ngram'
show(object)
```

**Arguments**

x, object      An ngram object.

output          a character string; determines what exactly is printed. Options are "summary", "truncated", and "full".

**Details**

If output=="summary", then just a simple representation of the n-gram object will be printed; for example, "An ngram object with 5 2-grams".

If output=="truncated", then the n-grams will be printed up to a maximum of 5 total.

If output=="full" then all n-grams will be printed.

**See Also**

[ngram](#), [babble](#)



---

phrasetable	<i>Get Phrasetable</i>
-------------	------------------------

---

**Description**

Get a table

**Usage**

```
get.phrasetable(ng)
```

**Arguments**

ng	An ngram object.
----	------------------

**See Also**

[ngram-class](#)

**Examples**

```
library(ngram)

str <- "A B A C A B B"
ng <- ngram(str)
get.phrasetable(ng)
```

---

preprocess	<i>Basic Text Preprocessor</i>
------------	--------------------------------

---

**Description**

A simple text preprocessor for use with the `ngram()` function.

**Usage**

```
preprocess(x, case = "lower", remove.punct = FALSE,
  remove.numbers = FALSE, fix.spacing = TRUE)
```

**Arguments**

x	Input text.
case	Option to change the case of the text. Value should be "upper", "lower", or NULL (no change).
remove.punct	Logical; should punctuation be removed?
remove.numbers	Logical; should numbers be removed?
fix.spacing	Logical; should multi/trailing spaces be collapsed/removed.

**Details**

The input text `x` must already be in the correct form for `ngram()`, i.e., a single string (character vector of length 1).

The `case` argument can take 3 possible values: `NULL`, in which case nothing is done, or `lower` or `upper`, wherein the case of the input text will be made lower/upper case, respectively.

**Value**

`concat()` returns

**Examples**

```
library(ngram)

x <- "Watch out for snakes! 111"
preprocess(x)
preprocess(x, remove.punct=TRUE, remove.numbers=TRUE)
```

---

rcorpus

*Random Corpus*


---

**Description**

Generate a corpus of random "words".

**Usage**

```
rcorpus(nwords = 50, alphabet = letters, minwordlen = 1, maxwordlen = 6)
```

**Arguments**

<code>nwords</code>	Number of words to generate.
<code>alphabet</code>	The pool of "letters" that word generation coes from. By default, it is the lower-case roman alphabet.
<code>minwordlen</code> , <code>maxwordlen</code>	The min/max length of words in the generated corpus.

**Value**

A string.

**Examples**

```
rcorpus(10)
```

---

splitter*Character Splitter*

---

**Description**

A utility function for use with n-gram modeling. This function splits a string based on various options.

**Usage**

```
splitter(string, split.char = FALSE, split.space = TRUE, spacesep = "_",
         split.punct = FALSE)
```

**Arguments**

string	An input string.
split.char	Logical; should a split occur after every character?
split.space	Logical; determines if spaces should be preserved as characters in the n-gram tokenization. The character(s) used for spaces are determined by the spacesep argument. characters.
spacesep	The character(s) to represent a space in the case that split.space=TRUE. Should not just be a space(s).
split.punct	Logical; determines if splits should occur at punctuation.

**Details**

Note that choosing split.char=TRUE necessarily implies split.punct=TRUE as well — but *not* necessarily that split.space=TRUE.

**Value**

A string.

**Examples**

```
x <- "watch out! a snake!"

splitter(x, split.char=TRUE)
splitter(x, split.space=TRUE, spacesep="_")
splitter(x, split.punct=TRUE)
```

---

string.summary	<i>Text Summary</i>
----------------	---------------------

---

**Description**

Text Summary

**Usage**

```
string.summary(string, wordlen_max = 10, senlen_max = 10, syllen_max = 10)
```

**Arguments**

string            An input string.  
wordlen\_max, senlen\_max, syllen\_max  
                  The maximum lengths of words/sentences/syllables to consider.

**Value**

A list of class string\_summary.

**Examples**

```
x <- "a b a c a b b"  
  
string.summary(x)
```

---

Tokenize	<i>n-gram Tokenization</i>
----------	----------------------------

---

**Description**

The ngram() function is the main workhorse of this package. It takes an input string and converts it into the internal n-gram representation.

**Usage**

```
ngram(str, n = 2, sep = " ")  
  
## S4 method for signature 'character'  
ngram(str, n = 2, sep = " ")
```

**Arguments**

<code>str</code>	The input text.
<code>n</code>	The 'n' as in 'n-gram'.
<code>sep</code>	A set of separator characters for the "words". See details for information about how this works; it works a little differently from <code>sep</code> arguments in R functions.

**Details**

On evaluation, a copy of the input string is produced and stored as an external pointer. This is necessary because the internal list representation just points to the first char of each word in the input string. So if you (or R's gc) deletes the input string, basically all hell breaks loose.

The `sep` parameter splits at any of the characters in the string. So `sep=" , "` splits at a comma or a space.

**Value**

An ngram class object.

**See Also**

[ngram-class](#), [getters](#), [phrasetable](#), [babble](#)

**Examples**

```
library(ngram)

str <- "A B A C A B B"
ngram(str, n=2)

str <- "A,B,A,C A B B"
### Split at a space
print(ngram(str), output="full")
### Split at a comma
print(ngram(str, sep=","), output="full")
### Split at a space or a comma
print(ngram(str, sep=" , "), output="full")
```

**Description**

An n-gram tokenizer with identical output to the `NGramTokenizer` function from the `RWeka` package.

**Usage**

```
ngram_asweka(str, min = 2, max = 2, sep = " ")
```

**Arguments**

<code>str</code>	The input text.
<code>min, max</code>	The minimum and maximum 'n' as in 'n-gram'.
<code>sep</code>	A set of separator characters for the "words". See details for information about how this works; it works a little differently from <code>sep</code> arguments in R functions.

**Details**

This n-gram tokenizer behaves similarly in both input and return to the tokenizer in RWeka. Unlike the tokenizer `ngram()`, the return is not a special class of external pointers; it is a vector, and therefore can be serialized via `save()` or `saveRDS()`.

**Value**

A vector of n-grams listed in decreasing blocks of n, in order within a block. The output matches that of RWeka's n-gram tokenizer.

**See Also**

[ngram](#)

**Examples**

```
library(ngram)

str <- "A B A C A B B"
ngram_asweka(str, min=2, max=4)
```

---

wordcount

*wordcount*

---

**Description**

`wordcount()` counts words. Currently a "word" is a clustering of characters separated from another clustering of characters by at least 1 space. That is the law.

**Usage**

```
wordcount(x, sep = " ", count.function = sum)

## S4 method for signature 'character'
wordcount(x, sep = " ", count.function = sum)

## S4 method for signature 'ngram'
wordcount(x, sep = " ", count.function = sum)
```

**Arguments**

x	A string or vector of strings, or an ngram object.
sep	The characters used to separate words.
count.function	The function to use for aggregation.

**Value**

A count.

**See Also**

[preprocess](#)

**Examples**

```
library(ngram)

words <- c("a", "b", "c")
words
wordcount(words)

str <- concatenate(words, collapse="")
str
wordcount(str)
```

# Index

- \*Topic **Amusement**
    - [babble](#), [3](#)
  - \*Topic **Package**
    - [ngram-package](#), [2](#)
  - \*Topic **Preprocessing**
    - [concatenate](#), [4](#)
    - [preprocess](#), [9](#)
    - [splitter](#), [11](#)
  - \*Topic **Summarize**
    - [string.summary](#), [12](#)
    - [wordcount](#), [14](#)
  - \*Topic **Tokenization**
    - [getters](#), [5](#)
    - [ngram-class](#), [7](#)
    - [phrasetable](#), [9](#)
    - [Tokenize](#), [12](#)
    - [Tokenize-AsWeka](#), [13](#)
  - \*Topic **Utility**
    - [getseed](#), [5](#)
    - [multiread](#), [6](#)
    - [rcorpus](#), [10](#)
- [babble](#), [3](#), [5](#), [8](#), [13](#)  
[babble](#), [ngram-method \(babble\)](#), [3](#)
- [concatenate](#), [4](#)
- [get.nextwords \(getters\)](#), [5](#)  
[get.nextwords](#), [ngram-method \(getters\)](#), [5](#)  
[get.ngrams \(getters\)](#), [5](#)  
[get.ngrams](#), [ngram-method \(getters\)](#), [5](#)  
[get.phrasetable \(phrasetable\)](#), [9](#)  
[get.string \(getters\)](#), [5](#)  
[get.string](#), [ngram-method \(getters\)](#), [5](#)  
[getseed](#), [3](#), [5](#)  
[getters](#), [5](#), [13](#)
- [multiread](#), [6](#)
- [ngram](#), [3](#), [6–8](#), [14](#)  
[ngram \(Tokenize\)](#), [12](#)
- [ngram](#), [character-method \(Tokenize\)](#), [12](#)  
[ngram-class](#), [7](#)  
[ngram-package](#), [2](#)  
[ngram-print](#), [8](#)  
[ngram\\_asweka \(Tokenize-AsWeka\)](#), [13](#)
- [phrasetable](#), [9](#), [13](#)  
[preprocess](#), [4](#), [9](#), [15](#)  
[print](#), [ngram-method \(ngram-print\)](#), [8](#)
- [rcorpus](#), [10](#)
- [show](#), [ngram-method \(ngram-print\)](#), [8](#)  
[splitter](#), [11](#)  
[string.summary](#), [12](#)
- [Tokenize](#), [8](#), [12](#)  
[tokenize \(Tokenize\)](#), [12](#)  
[Tokenize-AsWeka](#), [13](#)
- [wordcount](#), [14](#)  
[wordcount](#), [character-method \(wordcount\)](#), [14](#)  
[wordcount](#), [ngram-method \(wordcount\)](#), [14](#)